


Requirements for a Healthy Ecosystem in Advertising (RHEA)

Building next-generation, pro-democracy advertising
technology

Robin Berjon

2025-08-08



supramundane *agency.*

Advertising is a critical input to much of the digital economy, not least of which media. In fact, it has proven essential to the point of being described as “the business model of the internet.” Unfortunately, however, advertising has not been governed with the care that such an important piece of infrastructure deserves. Advertising is unpopular with technologists who want to do good, activists and civil society organisations often simply propose to eliminate it or limit themselves to superficial proposals that don’t engage with the technology or economics (e.g. “just use contextual”), and regulators have found the advertising system highly opaque and complex to engage with. It is common that publishers and marketers would only partially understand the system they use, which in turn stokes fear of what may happen if anything were to change. The field is highly concentrated around a Google/Meta duopoly and fraught with a long list of ills, but a comprehensive vision for a stakeholder-centric alternative has yet to emerge.

This document outlines what an alternative could be. It doesn’t claim to be complete and certainly not detailed, but it offers a path forward that can hopefully be iterated upon and deployed.

An Ecosystem Of Multifaceted Problems

There is extensive literature about issues in digital advertising, and summarising it is beyond the scope of this paper. Instead, we offer a list of the more salient problems in the space so as to inform both the urgent need for reform and the shape of the solution space we consider further down.

Taxing The Media to Subsidise Disinformation

Legacy digital advertising is a system that, at the scale of the entire internet, taxes high-quality contexts (often the media) and redistributes that money to slop farms and disinformation operators, while advertising intermediaries take a cut.

To understand how that works, it is important to understand that today's digital advertising shares data in such a way that the same person will be recognised by intermediaries across completely unrelated contexts. Crucially, this means that a person can be observed in one context and targeted in another.

Not all contexts are equal. Quality publishers (of all sizes) develop specific

audiences over time and through hard work. These audiences are valuable for advertisers because they correspond to a meaningful group of people that can correlate well with brand goals or a given purchase intent. Slop farms and disinformation networks, on the other hand, do not produce valuable, well-developed audiences. Instead, they try to drive traffic up through any means available. In addition to the low quality of the context they offer, they also have amorphous audiences without clear advertising value. When contexts are isolated (and also when advertisers have better control over where their advertising ends up), it is more challenging for slop and disinformation to be profitable.

This all changes when people can be targeted across contexts. A given person can be observed as being part of a high-value audience and later targeted in a low-value context. This both decreases the valuation of the quality context — since it is getting fewer, lower bids for its own audience — and increases the revenue that flows to slop and disinformation, what is



known as *Made For Advertising (MFA)* sites. It is a distributional arrangement that takes money from hard-working media to subsidise disinformation. This is detrimental to the media, to democracy, and to the information ecosystem overall.

A few additional details add some colour. First, slop is profitable on volume (of which there is a lot) and also provides an outlet for lower-quality ads. There is an equilibrium point that blends a given proportion of quality and slop contexts at which an advertising marketplace such as Google's is maximally profitable. Since the marketplace controls data distribution and prices, they are in a position to ensure this equilibrium is maintained.

Second, it's tempting to believe that marketers may benefit from lower prices thanks to this but it's not obvious that they do — they often have poor visibility into the contexts in which their advertising reaches their intended clients.

And finally, it's important to keep in mind that what policy documents from Google or Meta describe as “publishers” is a large undifferentiated group that almost always includes active disinformation networks and slop farms.

A pro-democracy solution must put an end to this subsidy system, both to make quality publishers more profitable and disinformation a less attractive business.

Privacy

tk

- privacy
 - limited evidence to support that the data processing is needed
 - often used to predict a sale that would have happened anyway (HBR)
 - the data advantage is most likely about faking lift than it is

about creating genuine conversions

- very little academic research worth considering, in some cases preposterously wrong
 - so opaque that no evidence-based decision-making is possible
 - the general climate of opacity makes it impossible for actors to make informed decisions as market participants, either as buyers or as sellers
- price
- advertiser privacy
- taxation
- money traceability
- the excessive weight of search & the browser funding system (put this at end)
- marketplace capture
- energy costs (for fraud, but also for oversized creatives)

Fraud

tk

Security

tk

Espionage

tk

Opacity

tk

Concentration and Governance

tk

Sustainability

tk

Our Goals: Who To Improve Advertising For?

tk

- ☐ Better for publishers: describe the subsidy system and how we need to end it
- ☐ Better for advertisers: traceability problems, lots of marketing options are bad
- ☐ Better for people: credible privacy
- ☐ Governed ad safety (link to defunding climate reporting + stealing content)
- ☐ from RHEA
- ☐ Exemptions + technical guarantees for the processing
- ☐ Explain that this is intended to help publishers too
- ☐ Single controller operations (apart perhaps from some very limited technical services like serving)
- ☐ Article 40/41

- ☐ Could consider using the Data Act and data exchanges (?) to support some form of targeting
 - ☐ Use sortition-based privacy (see Asia, also Max von G.)
- ☐ Beckn specialised for this
- ☐ SDAs — why this makes them *actually* possible
- ☐ Some geo
- ☐ PPA
- ☐ Fraud prevention
- ☐ KYC requirements
- ☐ Money: can this use a digital currency and protocol that supports splitting well?
- ☐ Locked down, but that's okay because purchases are made in small batches
- ☐



Content-addressed and all content is available at purchase time (this allows extra charges for creative size to be addressed)

- ☐ Bring back some GARUDA
- ☐

I don't think that we want browsers involved

- ☐ Split the components of the architecture so that we can have different governance for different elements — AT-style

Privacy

tk

Marketplace

tk

Creatives

tk

Governance

tk